

# Kernel Methods in Machine Learning

Julien Mairal & Jean-Philippe Vert

Romain Ménégaux (T.A.)



`firstname.lastname@m4x.org`



Google AI



# General learning framework





## Input

- $\mathcal{X}$  the space of **patterns** or **data** (typically,  $\mathcal{X} = \mathbb{R}^p$ )
- $\mathcal{Y}$  the space of **response** or **labels**
  - Classification or pattern recognition :  $\mathcal{Y} = \{-1, 1\}$
  - Regression :  $\mathcal{Y} = \mathbb{R}$
  - Structured output:  $\mathcal{Y}$  general
- $\mathcal{S} = \{(x_1, y_1), \dots, (x_n, y_n)\}$  a **training set** in  $(\mathcal{X} \times \mathcal{Y})^n$

## Output

- A **function**  $f : \mathcal{X} \rightarrow \mathcal{Y}$  to predict the output associated to any new pattern  $x \in \mathcal{X}$  by  $f(x)$

# What's wrong?



- OLS: the linear separation is not appropriate = "large bias"
- 1-NN: the classifier seems too unstable = "large variance"

# The fundamental "bias-variance" trade-off

- Assume  $Y = f(X) + \epsilon$ , where  $\epsilon$  is some noise
- From the training set  $\mathcal{S}$  we estimate the predictor  $\hat{f}$
- On a new point  $x_0$ , we predict  $\hat{f}(x_0)$  but the "true" observation will be  $Y_0 = f(x_0) + \epsilon$
- On average, we make an error of:

$$\begin{aligned} E_{\epsilon, \mathcal{S}} \left( Y_0 - \hat{f}(x_0) \right)^2 \\ &= E_{\epsilon, \mathcal{S}} \left( f(x_0) + \epsilon - \hat{f}(x_0) \right)^2 \\ &= E_{\epsilon}^2 + E_{\mathcal{S}} \left( f(x_0) - \hat{f}(x_0) \right)^2 \\ &= E_{\epsilon}^2 + \left( f(x_0) - E_{\mathcal{S}} \hat{f}(x_0) \right)^2 + E_{\mathcal{S}} \left( \hat{f}(x_0) - E_{\mathcal{S}} \hat{f}(x_0) \right)^2 \\ &= \text{noise} + \text{bias}^2 + \text{variance} \end{aligned}$$

# Important message

$$\text{Future prediction error} = \text{noise} + \text{bias}^2 + \text{variance}$$

- The "noise" part can not be avoided
- By choosing a learning model, we should consider both "bias" and "variance" if we want to make good predictions
- Intuitively, a **more realistic, more complex model** with more parameters to estimate has **smaller bias** but **larger variance**
- If variance dominates bias (eg, in high dimension), then having more complex, more realist models can hurt performance
- In other words, **a wrong but simple model can work better than a more realistic but more complex model**
- In many applications, domain experts (non-statisticians) often ignore the cost of complexity and prefer complex models, which can lead to disappointing results. You can help them!

- Linear model with parameter  $\beta \in \mathbb{R}^p$ :

$$\forall \mathbf{x} \in \mathbb{R}^p, \quad f_{\beta}(\mathbf{x}) = \beta^{\top} \mathbf{x} \quad \left( = \sum_{i=1}^p \beta_i x_i \right)$$

- Estimate  $\hat{\beta}^{OLS}$  from training data to minimize the mean sum of squares (MSE):

$$\text{MSE}(\beta) = \frac{1}{n} \sum_{i=1}^n (y_i - f_{\beta}(x_i))^2$$

## Back to OLS (cont.)

- Let's use matrix notations:

- $Y = (y_1, \dots, y_n)^\top \in \mathbb{R}^n$  the vector of outcomes
- $X = (x_1, \dots, x_n)^\top \in \mathbb{R}^{n \times p}$  the matrix ( $n$  rows=samples,  $p$  columns=features)

- We can rewrite MSE as

$$\text{MSE}(\beta) = \frac{1}{n} (Y - X\beta)^\top (Y - X\beta)$$

- $\text{MSE}(\beta)$  is a quadratic convex function; we minimize it by setting its gradient to 0:

$$\nabla_{\beta} \text{MSE}(\beta) = \frac{2}{n} X^\top (X\beta - Y) = 0$$

- If  $X^\top X$  is non-singular, the minimum is reached at

$$\hat{\beta}^{OLS} = \underset{\beta}{\operatorname{argmin}} \text{MSE}(\beta) = (X^\top X)^{-1} X^\top Y$$

# Properties of OLS

## Bias and variance of OLS

- Assume  $Y = X\beta^* + \epsilon$ , where  $E\epsilon = 0$  and  $E\epsilon\epsilon^\top = \sigma^2 I$ .
- Then the least squares estimator

$$\hat{\beta}^{OLS} = (X^\top X)^{-1} X^\top Y$$

satisfies

$$\begin{cases} E(\hat{\beta}^{OLS}) = \beta^*, \\ \text{Var}(\hat{\beta}^{OLS}) = E(\hat{\beta}^{OLS} - \beta^*)(\hat{\beta}^{OLS} - \beta^*)^\top = \sigma^2 (X^\top X)^{-1}. \end{cases}$$

Proof: exercise



# A solution: shrinkage estimators

- 1 Define a large family of "candidate classifiers", e.g., **linear predictors**:

$$f_{\beta}(x) = \beta^{\top} x \quad \text{for } x \in \mathbb{R}^p$$

# A solution: shrinkage estimators

- 1 Define a large family of "candidate classifiers", e.g., **linear predictors**:

$$f_{\beta}(x) = \beta^{\top} x \quad \text{for } x \in \mathbb{R}^p$$

- 2 For any candidate classifier  $f_{\beta}$ , quantify how "good" it is on the training set with some **empirical risk**, e.g.:

$$R(\beta) = \frac{1}{n} \sum_{i=1}^n (f_{\beta}(x_i) - y_i)^2.$$

# A solution: shrinkage estimators

- 1 Define a large family of "candidate classifiers", e.g., **linear predictors**:

$$f_{\beta}(x) = \beta^{\top} x \quad \text{for } x \in \mathbb{R}^p$$

- 2 For any candidate classifier  $f_{\beta}$ , quantify how "good" it is on the training set with some **empirical risk**, e.g.:

$$R(\beta) = \frac{1}{n} \sum_{i=1}^n (f_{\beta}(x_i) - y_i)^2.$$

- 3 Choose  $\beta$  that achieves the minimum empirical risk, subject to some **constraint**:

$$\min_{\beta} R(\beta) \quad \text{subject to} \quad \Omega(\beta) \leq C,$$

for some **penalty function**  $\Omega : \mathbb{R}^p \rightarrow \mathbb{R}^+$  and  $C \geq 0$ .

# Equivalent formulation

$$\min_{\beta} R(\beta) \quad \text{subject to} \quad \Omega(\beta) \leq C$$

is equivalent to

$$\min_{\beta} R(\beta) + \lambda \Omega(\beta)$$


- There exists a (not necessarily unique) correspondence between  $C$  and  $\lambda$  such that the solutions to both problems are the same.
- If  $C$  increase,  $\lambda$  decreases
- The formulation with  $\lambda$  is often preferred to implement the algorithm
- Proof: using Lagrangian duality (only true under some assumptions, eg,  $R$  and  $\Omega$  convex + Slater conditions, see later)

# Choice of $C$ or $\lambda$

- Choose a grid of values  $\Lambda$  for  $\lambda$  (or  $C$ )
- For each  $\lambda \in \Lambda$  (or  $C$ ) estimate the best model

$$\hat{\beta}_{\lambda} \in \operatorname{argmin}_{\beta} R(\beta) + \lambda \Omega(\beta)$$

- Select  $\hat{\beta} = \hat{\beta}_{\hat{\lambda}}$  to **minimize the bias-variance tradeoff**.



# Cross-validation

A simple and systematic procedure to estimate the risk (and to optimize the model's parameters)

- 1 Randomly divide the training set (of size  $n$ ) into  $K$  (almost) equal portions, each of size  $K/n$
- 2 For each portion, fit the model with different parameters on the  $K - 1$  other groups and test its performance on the left-out group
- 3 Average performance over the  $K$  groups, and take the parameter with the smallest average performance.

Taking  $K = 5$  or  $10$  is recommended as a good default choice.

# Summary

- 1 Many problems in modern machine learning involve models with many parameters (i.e., high dimension)
- 2 The total prediction error of a learning system is the sum of a **bias** and a **variance** error
- 3 In **high dimension**, the **variance** term often dominates
- 4 **Shrinkage methods** allow us to control the bias/variance trade-off
- 5 The choice of the **penalty** is where we can put **prior knowledge** to decrease bias
- 6 The parameter to control the bias-variance trade-off ( $C$  or  $\lambda$ ) is typically chosen by **cross-validation**, to minimize the test error.

# Outline


- 1 Learning in high dimension
- 2 Learning with  $\ell_2$  regularization
  - Ridge regression
  - Ridge logistic regression
  - Linear hard-margin SVM
  - Interlude: quick notes on constrained optimization
  - Back to hard-margin SVM
  - Soft-margin SVM
  - Large-margin classifiers
- 3 Learning with kernels
  - Kernel methods
  - Positive definite kernels and RKHS
  - Kernel examples
  - Multiple Kernel Learning (MKL)
- 4 Conclusion



# Overview

- We focus on a simple penalty function: the squared Euclidean norm

$$\Omega(\beta) = \|\beta\|^2 \quad \left( = \beta^\top \beta = \sum_{i=1}^p \beta_i^2 \right)$$



- This will allow us to derive many state-of-the-art linear methods:
  - Ridge regression
  - Ridge logistic regression
  - SVM and large-margin classifiers
- This will allow us to extend these linear methods to nonlinear models, using kernels

# Outline

- 1 Learning in high dimension
- 2 Learning with  $\ell_2$  regularization
  - Ridge regression
  - Ridge logistic regression
  - Linear hard-margin SVM
  - Interlude: quick notes on constrained optimization
  - Back to hard-margin SVM
  - Soft-margin SVM
  - Large-margin classifiers
- 3 Learning with kernels
  - Kernel methods
  - Positive definite kernels and RKHS
  - Kernel examples
  - Multiple Kernel Learning (MKL)
- 4 Conclusion

# Ridge regression (?)

- 1 Consider the set of **linear predictors**:

$$\forall \beta \in \mathbb{R}^p, \quad f_\beta(x) = \beta^\top x \quad \text{for } x \in \mathbb{R}^p.$$

# Ridge regression (?)

- 1 Consider the set of **linear predictors**:

$$\forall \beta \in \mathbb{R}^p, \quad f_\beta(x) = \beta^\top x \quad \text{for } x \in \mathbb{R}^p.$$

- 2 Consider the MSE as empirical risk:

$$R(\beta) = \frac{1}{n} \sum_{i=1}^n (f_\beta(x_i) - y_i)^2.$$

# Ridge regression (?)

- ① Consider the set of **linear predictors**:

$$\forall \beta \in \mathbb{R}^p, \quad f_\beta(x) = \beta^\top x \quad \text{for } x \in \mathbb{R}^p.$$

- ② Consider the MSE as empirical risk:

$$R(\beta) = \frac{1}{n} \sum_{i=1}^n (f_\beta(x_i) - y_i)^2.$$

- ③ Consider the **squared Euclidean norm** as a penalty:

$$\Omega(\beta) = \|\beta\|^2.$$

- The penalized risk can be written in matrix form:

$$\begin{aligned} R(\beta) + \lambda \Omega(\beta) &= \frac{1}{n} \sum_{i=1}^n (f_{\beta}(x_i) - y_i)^2 + \lambda \sum_{i=1}^p \beta_i^2 \\ &= \frac{1}{n} (Y - X\beta)^{\top} (Y - X\beta) + \lambda \beta^{\top} \beta. \end{aligned}$$

- The penalized risk can be written in matrix form:

$$\begin{aligned} R(\beta) + \lambda \Omega(\beta) &= \frac{1}{n} \sum_{i=1}^n (f_{\beta}(x_i) - y_i)^2 + \lambda \sum_{i=1}^p \beta_i^2 \\ &= \frac{1}{n} (Y - X\beta)^{\top} (Y - X\beta) + \lambda \beta^{\top} \beta. \end{aligned}$$

- Unique minimizer (by setting the gradient to 0):

$$\hat{\beta}_{\lambda}^{\text{ridge}} = \arg \min_{\beta \in \mathbb{R}^p} \{R(\beta) + \lambda \Omega(\beta)\} = \left( X^{\top} X + \lambda n I \right)^{-1} X^{\top} Y.$$

# Performance of ridge regression

## Lemma

Assume that:

- $Y = X\beta^* + \epsilon$ , where  $E\epsilon = 0$  and  $E\epsilon\epsilon^\top = \sigma^2 I$ .
- $X^\top X = nI_p$  (orthogonal design)

Then:

$$\begin{cases} \text{bias}(\hat{\beta}_\lambda^{\text{ridge}}) = E(\hat{\beta}_\lambda^{\text{ridge}}) - \beta^* = -\frac{\lambda}{1+\lambda}\beta^*, \\ \text{Var}(\hat{\beta}_\lambda^{\text{ridge}}) = \frac{\sigma^2}{n(1+\lambda)^2} I_p = \frac{1}{(1+\lambda)^2} \text{Var}(\hat{\beta}^{\text{OLS}}). \end{cases}$$

Proof: exercise



# Performance of ridge regression

## Corollary

For any  $\lambda \geq 0$  let

$$f(\lambda) = E_{\mathcal{S}, x_0} \left[ \text{bias}^2 \left( x_0^\top \hat{\beta}_\lambda^{\text{ridge}} \right) + \text{Var} \left( x_0^\top \hat{\beta}_\lambda^{\text{ridge}} \right) \right]$$

where  $E x_0 = 0$ ,  $E x_0 x_0^\top = I_p$ . Then  $f(\lambda)$  is minimum for

$$\lambda^* = \frac{\sigma^2 p}{n \|\beta^*\|^2}$$

and

$$f(\lambda^*) = \frac{f(0)f(+\infty)}{f(0) + f(+\infty)} \leq \min \{f(0), f(\infty)\}$$

where

$$f(0) = \sigma^2 p / n, \quad f(\infty) = \|\beta^*\|^2$$

Proof: exercise

$$\hat{\beta}_{\lambda}^{\text{ridge}} = \left( X^{\top} X + \lambda n I \right)^{-1} X^{\top} Y$$

### Corollary

- As  $\lambda \rightarrow 0$ ,  $\hat{\beta}_{\lambda}^{\text{ridge}} \rightarrow \hat{\beta}^{\text{OLS}}$  (low bias, high variance).
- As  $\lambda \rightarrow +\infty$ ,  $\hat{\beta}_{\lambda}^{\text{ridge}} \rightarrow 0$  (high bias, low variance).

# Generalization: $\ell_2$ -regularized learning

- A general  $\ell_2$ -penalized estimator is of the form

$$\min_{\beta} \{ R(\beta) + \lambda \| \beta \|^2 \} , \quad (1)$$

where

$$R(\beta) = \frac{1}{n} \sum_{i=1}^n \ell(f_{\beta}(x_i), y_i)$$

for some general loss functions  $\ell$ .



- Ridge regression corresponds to the particular loss

$$\ell(u, y) = (u - y)^2 .$$

- For general, **convex** losses, the problem (??) is strictly convex and has a **unique global minimum**, which can usually be found by **numerical algorithms** for convex optimization.

# Losses for regression

- Square loss :  $\ell(u, y) = (u - y)^2$
- $\epsilon$ -insensitive loss :  $\ell(u, y) = (|u - y| - \epsilon)_+$
- Huber loss : mixed quadratic/linear




# Outline

- 1 Learning in high dimension
- 2 Learning with  $\ell_2$  regularization
  - Ridge regression
  - Ridge logistic regression
  - Linear hard-margin SVM
  - Interlude: quick notes on constrained optimization
  - Back to hard-margin SVM
  - Soft-margin SVM
  - Large-margin classifiers
- 3 Learning with kernels
  - Kernel methods
  - Positive definite kernels and RKHS
  - Kernel examples
  - Multiple Kernel Learning (MKL)
- 4 Conclusion

# Binary classification

## Setting

- $\mathcal{X} = \mathbb{R}^p$  set of inputs
- $\mathcal{Y} = \{-1, 1\}$  binary outputs
- $\mathcal{S} = \{(x_1, y_1), \dots, (x_n, y_n)\}$  a training set in  $(\mathcal{X} \times \mathcal{Y})^n$
- Goal: Estimate a function  $f : \mathcal{X} \rightarrow \mathbb{R}$  to **predict  $y$  by  $\text{sign}(f(x))$**



# The 0/1 loss

- The 0/1 loss measures if a prediction is correct or not:

$$\ell_{0/1}(f(x), y) = \mathbf{1}(yf(x) < 0) = \begin{cases} 0 & \text{if } y = \text{sign}(f(x)) \\ 1 & \text{otherwise.} \end{cases}$$

- It is then tempting to learn  $f_\beta(x) = \beta^\top x$  by solving:


$$\min_{\beta \in \mathbb{R}^p} \underbrace{\frac{1}{n} \sum_{i=1}^n \ell_{0/1}(f_\beta(x_i), y_i)}_{\text{misclassification rate}} + \underbrace{\lambda \|\beta\|^2}_{\text{regularization}}$$

- However:
  - The problem is non-smooth, and typically NP-hard to solve
  - The regularization has **no effect** since the 0/1 loss is invariant by scaling of  $\beta$
  - In fact, no function achieves the minimum when  $\lambda > 0$  (why?)

# The logistic loss

- An alternative is to define a probabilistic model of  $y$  parametrized by  $f(x)$ , e.g.:

$$\forall y \in \{-1, 1\}, \quad p(y | f(x)) = \frac{1}{1 + e^{-yf(x)}} = \sigma(yf(x))$$



- The **logistic loss** is the negative conditional likelihood:

$$\ell_{\text{logistic}}(f(x), y) = -\ln p(y | f(x)) = \ln(1 + e^{-yf(x)})$$



# Ridge logistic regression (?)

$$\min_{\beta \in \mathbb{R}^p} J(\beta) = \frac{1}{n} \sum_{i=1}^n \ln \left( 1 + e^{-y_i \beta^\top x_i} \right) + \lambda \|\beta\|^2$$

- Can be interpreted as a regularized conditional maximum likelihood estimator
- No explicit solution, but smooth convex optimization problem that can be solved numerically

## Solving ridge logistic regression

$$\min_{\beta} J(\beta) = \frac{1}{n} \sum_{i=1}^n \ln \left( 1 + e^{-y_i \beta^\top x_i} \right) + \lambda \|\beta\|^2$$

No explicit solution, but convex problem with:

$$\begin{aligned} \nabla_{\beta} J(\beta) &= -\frac{1}{n} \sum_{i=1}^n \frac{y_i x_i}{1 + e^{y_i \beta^\top x_i}} + 2\lambda \beta \\ &= -\frac{1}{n} \sum_{i=1}^n y_i [1 - P_{\beta}(y_i | x_i)] x_i + 2\lambda \beta \\ \nabla_{\beta}^2 J(\beta) &= \frac{1}{n} \sum_{i=1}^n \frac{x_i x_i^\top e^{y_i \beta^\top x_i}}{(1 + e^{y_i \beta^\top x_i})^2} + 2\lambda I \\ &= \frac{1}{n} \sum_{i=1}^n P_{\beta}(1 | x_i) (1 - P_{\beta}(1 | x_i)) x_i x_i^\top + 2\lambda I \end{aligned}$$

# Outline

## 1 Learning in high dimension

## 2 Learning with $\ell_2$ regularization

- Ridge regression
- Ridge logistic regression
- Linear hard-margin SVM
- Interlude: quick notes on constrained optimization
- Back to hard-margin SVM
- Soft-margin SVM
- Large-margin classifiers

## 3 Learning with kernels

- Kernel methods
- Positive definite kernels and RKHS
- Kernel examples
- Multiple Kernel Learning (MKL)

## 4 Conclusion

# Loss functions for classifications

We already saw 3 loss functions for binary classification problems

- The 0/1 loss  $\ell_{0/1}(f(x), y) = \mathbf{1}(yf(x) < 0)$
- The logistic loss  $\ell_{\text{logistic}}(f(x), y) = \ln(1 + e^{-yf(x)})$
- The hinge loss  $\ell_{\text{hinge}}(f(x), y) = \max(0, 1 - yf(x))$

## Definition

In binary classification ( $\mathcal{Y} = \{-1, 1\}$ ), the **margin** of the function  $f$  for a pair  $(x, y)$  is:


$$yf(x).$$

In all cases the loss is a decreasing function of the margin, i.e.,

$$\ell(f(x), y) = \varphi(yf(x)) \text{ , with } \varphi \text{ non-increasing}$$


What about other similar loss functions?

# Loss function examples





Method	$\varphi(u)$
Logistic regression	$\log(1 + e^{-u})$
Support vector machine (1-SVM)	$\max(1 - u, 0)$
Support vector machine (2-SVM)	$\max(1 - u, 0)^2$
Boosting	$e^{-u}$

## Summary: large margin classifiers



- $\varphi$  calibrated (e.g., decreasing,  $\varphi'(0) < 0$ )  $\implies$  good proxy for classification error
- $\varphi$  convex + representer theorem  $\implies$  efficient algorithms
- $\varphi$  smooth (Lipschitz) +  $\ell_2$  regularization  $\implies$  good learning ability

# Summary: $\ell_2$ -regularized linear methods



$$f_{\beta}(x) = \beta^{\top} x, \quad \min_{\beta} \left\{ \frac{1}{n} \sum_{i=1}^n \ell(f_{\beta}(x_i), y_i) + \lambda \|\beta\|_2^2 \right\}$$


- Many popular methods for regression and classification are obtained by changing the loss function: ridge regression, logistic regression, SVM...
- Needs to solve numerically a convex optimization problem, well adapted to large datasets (stochastic gradient...)
- In practice, very similar performance between the different variants in general

# Outline

- 1 Learning in high dimension
- 2 Learning with  $\ell_2$  regularization
  - Ridge regression
  - Ridge logistic regression
  - Linear hard-margin SVM
  - Interlude: quick notes on constrained optimization
  - Back to hard-margin SVM
  - Soft-margin SVM
  - Large-margin classifiers
- 3 Learning with kernels
  - Kernel methods
  - Positive definite kernels and RKHS
  - Kernel examples
  - Multiple Kernel Learning (MKL)
- 4 Conclusion



# Motivation



- Sometimes linear models are not interesting...
- Kernels will allow to solve nonlinear problems with linear methods!

# Outline

## 1 Learning in high dimension

## 2 Learning with $\ell_2$ regularization


- Ridge regression
- Ridge logistic regression
- Linear hard-margin SVM
- Interlude: quick notes on constrained optimization
- Back to hard-margin SVM
- Soft-margin SVM
- Large-margin classifiers

## 3 Learning with kernels

- **Kernel methods**
- Positive definite kernels and RKHS
- Kernel examples
- Multiple Kernel Learning (MKL)

## 4 Conclusion

"Linear" depends on the representation you choose



For  $x = \begin{pmatrix} x_1 \\ x_2 \end{pmatrix}$  let  $\Phi(x) = \begin{pmatrix} x_1^2 \\ x_2^2 \end{pmatrix}$ . The decision function is:

$$f(x) = x_1^2 + x_2^2 - R^2 = \beta^\top \Phi(x) + b$$

with  $\beta = (1, 1)^\top$  and  $b = -R^2$

# Kernel = inner product in the feature space


## Definition

For a given mapping


$$\Phi : \mathcal{X} \mapsto \mathcal{H}$$

from the space of data  $\mathcal{X}$  to some feature space  $\mathcal{H}$ , the **kernel** between two objects  $x$  and  $x'$  is the inner product of their images:

$$\forall x, x' \in \mathcal{X}, \quad K(x, x') = \Phi(x)^\top \Phi(x').$$



# Example




Let  $\mathcal{X} = \mathcal{H} = \mathbb{R}^2$  and for  $x = \begin{pmatrix} x_1 \\ x_2 \end{pmatrix}$  let  $\Phi(x) = \begin{pmatrix} x_1^2 \\ x_2^2 \end{pmatrix}$

Then the kernel is:

$$K(x, x') = \Phi(x)^\top \Phi(x') = (x_1)^2 (x'_1)^2 + (x_2)^2 (x'_2)^2.$$


# The kernel tricks



## 2 tricks

- 1 Many linear algorithms (in particular  $\ell_2$ -regularized methods) can be performed in the feature space of  $\Phi(x)$  **without explicitly computing the images  $\Phi(x)$ , but instead by computing kernels  $K(x, x')$ .**
- 2 It is sometimes possible to **easily** compute kernels which correspond to complex large-dimensional feature spaces:  **$K(x, x')$  is often much simpler to compute than  $\Phi(x)$  and  $\Phi(x')$**


## Trick 2 illustration: polynomial kernel



For  $x = (x_1, x_2)^\top \in \mathbb{R}^2$ , let  $\Phi(x) = (x_1^2, \sqrt{2}x_1x_2, x_2^2) \in \mathbb{R}^3$ :

$$\begin{aligned} K(x, x') &= x_1^2 x_1'^2 + 2x_1x_2x_1'x_2' + x_2^2 x_2'^2 \\ &= (x_1x_1' + x_2x_2')^2 \\ &= (x^\top x')^2. \end{aligned}$$

## Trick 2 illustration: polynomial kernel



More generally, for  $x, x' \in \mathbb{R}^p$ ,

$$K(x, x') = \left( x^\top x' + 1 \right)^d$$

is an inner product in a feature space of all monomials of degree up to  $d$   
(left as exercise.)



# More generally: trick 1 for $\ell_2$ -regularized linear models

## Representer theorem

Let  $f_\beta(x) = \beta^\top \Phi(x)$ . Then any solution  $\hat{f}_\beta$  of

$$\min_{\beta} \frac{1}{n} \sum_{i=1}^n \ell(f_\beta(x_i), y_i) + \lambda \|\beta\|_2^2$$

can be expanded as

$$\hat{f}_\beta(x) = \sum_{i=1}^n \alpha_i K(x_i, x),$$

where  $\alpha \in \mathbb{R}^n$  is a solution of:

$$\min_{\alpha \in \mathbb{R}^n} \frac{1}{n} \sum_{i=1}^n \ell \left( \sum_{j=1}^n \alpha_j K(x_i, x_j), y_i \right) + \lambda \sum_{i,j=1}^n \alpha_i \alpha_j K(x_i, x_j).$$

# Representer theorem: proof

- For any  $\beta \in \mathbb{R}^p$ , decompose  $\beta = \beta_S + \beta_\perp$  where  $\beta_S \in \text{span}(\Phi(x_1), \dots, \Phi(x_n))$  and  $\beta_\perp$  is orthogonal to it.
- On any point  $x_i$  of the training set, we have:

$$f_\beta(x_i) = \beta^\top \Phi(x_i) = \beta_S^\top \Phi(x_i) + \beta_\perp^\top \Phi(x_i) = \beta_S^\top \Phi(x_i) = f_{\beta_S}(x_i).$$

- On the other hand, we have  $\|\beta\|_2^2 = \|\beta_S\|_2^2 + \|\beta_\perp\|_2^2 \geq \|\beta_S\|_2^2$ , with strict inequality if  $\beta_\perp \neq 0$ .
- Consequently,  $\beta_S$  is always as good as  $\beta$  in terms of objective function, and strictly better if  $\beta_\perp \neq 0$ . This implies that at any minimum,  $\beta_\perp = 0$  and therefore  $\beta = \beta_S = \sum_{i=1}^n \alpha_i \Phi(x_i)$  for some  $\alpha \in \mathbb{R}^n$ .
- We then just replace  $\beta$  by this expression in the objective function, noting that

$$\|\beta\|_2^2 = \left\| \sum_{i=1}^n \alpha_i \Phi(x_i) \right\|_2^2 = \sum_{i,j=1}^n \alpha_i \alpha_j \Phi(x_i)^\top \Phi(x_j) = \sum_{i,j=1}^n \alpha_i \alpha_j K(x_i, x_j).$$

## Example: kernel ridge regression

- Let  $\Phi : \mathcal{X} \rightarrow \mathbb{R}^p$  be a feature mapping from the space of data to a Euclidean or Hilbert space.
- Let  $f_\beta(x) = \beta^\top \Phi(x)$  and  $K$  the corresponding kernel.
- By the representer theorem, any solution of:

$$\hat{f} = \arg \min_{f_\beta} \frac{1}{n} \sum_{i=1}^n (y_i - f_\beta(x_i))^2 + \lambda \|\beta\|_2^2$$

can be expanded as:

$$\hat{f} = \sum_{i=1}^n \alpha_i K(x_i, x).$$

## Example: kernel ridge regression

- Let  $Y = (y_1, \dots, y_n)^\top \in \mathbb{R}^n$  the vector of response variables.
- Let  $\alpha = (\alpha_1, \dots, \alpha_n)^\top \in \mathbb{R}^n$  the unknown coefficients.
- Let  $K$  be the  $n \times n$  Gram matrix:  $K_{i,j} = K(x_i, x_j)$ .
- We can then write in matrix form:

$$\left( \hat{f}(x_1), \dots, \hat{f}(x_n) \right)^\top = K\alpha,$$

- Moreover,

$$\|\beta\|_2^2 = \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j K(x_i, x_j) = \alpha^\top K \alpha.$$

## Example: kernel ridge regression

- The problem is therefore equivalent to:

$$\arg \min_{\alpha \in \mathbb{R}^n} \frac{1}{n} (K\alpha - Y)^\top (K\alpha - Y) + \lambda \alpha^\top K \alpha.$$


- This is a convex and differentiable function of  $\alpha$ . Its minimum can therefore be found by setting the gradient in  $\alpha$  to zero:

$$\begin{aligned} 0 &= \frac{2}{n} K (K\alpha - Y) + 2\lambda K \alpha \\ &= K [(K + \lambda nI) \alpha - Y] \end{aligned}$$


- For  $\lambda > 0$ ,  $K + \lambda nI$  is invertible (because  $K$  is positive semidefinite) so one solution is to take:

$$\alpha = (K + \lambda nI)^{-1} Y.$$


## Example (KRR with Gaussian RBF kernel)



# Example (KRR with Gaussian RBF kernel)




# Example (KRR with Gaussian RBF kernel)






# Example (KRR with Gaussian RBF kernel)

**lambda = 10**




# Example (KRR with Gaussian RBF kernel)




# Example (KRR with Gaussian RBF kernel)

**lambda = 0.1**




# Example (KRR with Gaussian RBF kernel)




# Example (KRR with Gaussian RBF kernel)


**lambda = 0.001**




# Example (KRR with Gaussian RBF kernel)



# Example (KRR with Gaussian RBF kernel)




# Example (KRR with Gaussian RBF kernel)





# Example (KRR with Gaussian RBF kernel)

**lambda = 0.0000001**



## Remark: uniqueness of the solution

Let us find *all*  $\alpha$ 's that solve

$$K [(K + \lambda nI) \alpha - Y] = 0$$

- $K$  being a symmetric matrix, it can be diagonalized in an orthonormal basis and  $\text{Ker}(K) \perp \text{Im}(K)$ .
- In this basis we see that  $(K + \lambda nI)^{-1}$  leaves  $\text{Im}(K)$  and  $\text{Ker}(K)$  invariant.
- The problem is therefore equivalent to:

$$\begin{aligned}(K + \lambda nI) \alpha - Y &\in \text{Ker}(K) \\ \Leftrightarrow \alpha - (K + \lambda nI)^{-1} Y &\in \text{Ker}(K) \\ \Leftrightarrow \alpha &= (K + \lambda nI)^{-1} Y + \epsilon, \text{ with } K\epsilon = 0.\end{aligned}$$

- However, if  $\alpha' = \alpha + \epsilon$  with  $K\epsilon = 0$ , then:

$$\|\beta - \beta'\|_2^2 = (\alpha - \alpha')^\top K (\alpha - \alpha') = 0,$$

therefore  $\beta = \beta'$ . **KRR has a unique solution  $\beta$ , which can possibly be expressed by several  $\alpha$ 's if  $K$  is singular.**

## Comparison with "standard" ridge regression

- Let  $X$  the  $n \times p$  data matrix,  $K = XX^\top$  the kernel Gram matrix.
- In "standard" ridge regression, we have  $\hat{f}(x) = \hat{\beta}^\top x$  with

$$\hat{\beta} = \left( X^\top X + n\lambda I \right)^{-1} X^\top Y.$$

- In "kernel" ridge regression, we have  $\tilde{f}(x) = \sum_{i=1}^n \alpha_i x_i^\top x = \tilde{\beta}^\top x$  with

$$\tilde{\beta} = \sum_{i=1}^n \alpha_i x_i = X^\top \alpha = X^\top \left( XX^\top + \lambda n I \right)^{-1} Y.$$

- Oups... which one is correct?

# Comparison with "standard" ridge regression

## Matrix inversion lemma

For any matrices  $B$  and  $C$ , and  $\gamma > 0$  the following holds (when it makes sense):

$$B (CB + \gamma I)^{-1} = (BC + \gamma I)^{-1} B$$

We deduce that (of course...):

$$\hat{\beta} = \underbrace{\left( X^T X + n\lambda I \right)^{-1}}_{p \times p} X^T Y = X^T \underbrace{\left( X X^T + \lambda n I \right)^{-1}}_{n \times n} Y = \tilde{\beta}$$

Computationally, inverting the matrix is the expensive part, which suggest to implement:

- KRR when  $p > n$  (high dimension)
- RR when  $p < n$  (many points)

# Generalization

- We learn the function  $f(x) = \sum_{i=1}^n \alpha_i K(x_i, x)$  by solving in  $\alpha$  the following optimization problem, with adequate loss function  $\ell$ :


$$\min_{\alpha \in \mathbb{R}^n} \frac{1}{n} \sum_{i=1}^n \ell \left( \sum_{j=1}^n \alpha_j K(x_i, x_j), y_i \right) + \lambda \sum_{i,j=1}^n \alpha_i \alpha_j K(x_i, x_j).$$

- No explicit solution, but **convex** optimization problem
- Note that the **dimension** of the problem is now  $n$  instead of  $p$  (useful when  $n < p$ )

# Outline

- 1 Learning in high dimension
- 2 Learning with  $\ell_2$  regularization
  - Ridge regression
  - Ridge logistic regression
  - Linear hard-margin SVM
  - Interlude: quick notes on constrained optimization
  - Back to hard-margin SVM
  - Soft-margin SVM
  - Large-margin classifiers
- 3 Learning with kernels
  - Kernel methods
  - Positive definite kernels and RKHS
  - Kernel examples
  - Multiple Kernel Learning (MKL)
- 4 Conclusion

## Remember: polynomial kernel



$$\forall x, x' \in \mathbb{R}^p, \quad K(x, x') = (x^\top x' + 1)^d$$

is an inner product in a feature space of all monomials of degree up to  $d$

# Which functions $K(x, x')$ are kernels?


## Definition

A function  $K(x, x')$  defined on a set  $\mathcal{X}$  is a **kernel** if and only if there exists a features space (Hilbert space)  $\mathcal{H}$  and a mapping

$$\Phi : \mathcal{X} \mapsto \mathcal{H} ,$$

such that, for any  $x, x'$  in  $\mathcal{X}$ :

$$K(x, x') = \langle \Phi(x), \Phi(x') \rangle_{\mathcal{H}} .$$





## Reminder ...

- An **inner product** on an  $\mathbb{R}$ -vector space  $\mathcal{H}$  is a mapping  $(f, g) \mapsto \langle f, g \rangle_{\mathcal{H}}$  from  $\mathcal{H}^2$  to  $\mathbb{R}$  that is **bilinear**, **symmetric** and such that  $\langle f, f \rangle > 0$  for all  $f \in \mathcal{H} \setminus \{0\}$ .
- A vector space endowed with an inner product is called **pre-Hilbert**. It is endowed with a norm defined by the inner product as
$$\|f\|_{\mathcal{H}} = \langle f, f \rangle_{\mathcal{H}}^{\frac{1}{2}}.$$
- A **Hilbert space** is a pre-Hilbert space **complete** for the norm defined by the inner product.

# Kernel examples

- **Polynomial** (on  $\mathbb{R}^d$ ):

$$K(x, x') = (x \cdot x' + 1)^d$$

- **Gaussian radial basis function (RBF)** (on  $\mathbb{R}^d$ )

$$K(x, x') = \exp\left(-\frac{\|x - x'\|^2}{2\sigma^2}\right)$$

- **Laplace** kernel (on  $\mathbb{R}$ )

$$K(x, x') = \exp(-\gamma|x - x'|)$$

- **Min** kernel (on  $\mathbb{R}_+$ )

$$K(x, x') = \min(x, x')$$

## Example: SVM with a Gaussian kernel

- Training:

$$\min_{\alpha \in \mathbb{R}^n} \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i,j=1}^n \alpha_i \alpha_j y_i y_j \exp \left( -\frac{\|\vec{x}_i - \vec{x}_j\|^2}{2\sigma^2} \right)$$

s.t.  $0 \leq \alpha_i \leq C$ , and  $\sum_{i=1}^n \alpha_i y_i = 0$ .


- Prediction

$$f(\vec{x}) = \sum_{i=1}^n \alpha_i \exp \left( -\frac{\|\vec{x} - \vec{x}_i\|^2}{2\sigma^2} \right)$$

# Example: SVM with a Gaussian kernel

$$f(\vec{x}) = \sum_{i=1}^n \alpha_i \exp\left(-\frac{\|\vec{x} - \vec{x}_i\|^2}{2\sigma^2}\right)$$

SVM classification plot



# Positive Definite (p.d.) functions

## Definition

A **positive definite (p.d.) function** on the set  $\mathcal{X}$  is a function  $K : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$  **symmetric**:

$$\forall (x, x') \in \mathcal{X}^2, \quad K(x, x') = K(x', x),$$


and which satisfies, for all  $N \in \mathbb{N}$ ,  $(x_1, x_2, \dots, x_N) \in \mathcal{X}^N$  et  $(a_1, a_2, \dots, a_N) \in \mathbb{R}^N$ :

$$\sum_{i=1}^N \sum_{j=1}^N a_i a_j K(x_i, x_j) \geq 0.$$

# Kernels are p.d. functions

Theorem (Aronszajn, 1950)

$K$  is a kernel *if and only if* it is a positive definite function.



# Proof: kernel $\implies$ p.d. (easy)

Let

$$K(x, x') = \langle \Phi(x), \Phi(x') \rangle_{\mathcal{H}}$$

be a kernel. It is p.d. because:

- $K(x, x') = \langle \Phi(x), \Phi(x') \rangle_{\mathcal{H}} = \langle \Phi(x'), \Phi(x) \rangle_{\mathcal{H}} = K(x', x)$  ,
- $\sum_{i=1}^N \sum_{j=1}^N a_i a_j \langle \Phi(x_i), \Phi(x_j) \rangle_{\mathcal{H}} = \| \sum_{i=1}^N a_i \Phi(x_i) \|^2_{\mathcal{H}} \geq 0$  .

## Proof: p.d. $\implies$ kernel when $\mathcal{X}$ is finite

- Suppose  $\mathcal{X} = \{x_1, x_2, \dots, x_N\}$  is finite of size  $N$ .
- Any p.d. kernel  $K : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$  is entirely defined by the  $N \times N$  symmetric positive semidefinite matrix  $[K]_{ij} := K(x_i, x_j)$ .
- It can therefore be diagonalized on an orthonormal basis of eigenvectors  $(u_1, u_2, \dots, u_N)$ , with non-negative eigenvalues  $0 \leq \lambda_1 \leq \dots \leq \lambda_N$ , i.e.,

$$K(x_i, x_j) = \left[ \sum_{l=1}^N \lambda_l u_l u_l^\top \right]_{ij} = \sum_{l=1}^N \lambda_l u_l(i) u_l(j) = \langle \Phi(x_i), \Phi(x_j) \rangle_{\mathbb{R}^N},$$

with

$$\Phi(x_i) = \begin{pmatrix} \sqrt{\lambda_1} u_1(i) \\ \vdots \\ \sqrt{\lambda_N} u_N(i) \end{pmatrix}. \quad \square$$



## Proof: p.d. $\implies$ kernel in the general case

- Mercer (1909) for  $\mathcal{X} = [a, b] \subset \mathbb{R}$  (more generally  $\mathcal{X}$  compact) and  $K$  continuous (the so-called **Mercer kernels**).
- Kolmogorov (1941) for  $\mathcal{X}$  countable.
- Aronszajn (1944, 1950) for the general case, using the theory of RKHS.

## Definition

Let  $\mathcal{X}$  be a set and  $\mathcal{H} \subset \mathbb{R}^{\mathcal{X}}$  be a **class of functions forming a (real) Hilbert space** with inner product  $\langle \cdot, \cdot \rangle_{\mathcal{H}}$ . The function  $K : \mathcal{X}^2 \mapsto \mathbb{R}$  is called a **reproducing kernel (r.k.)** of  $\mathcal{H}$  if

- ①  $\mathcal{H}$  contains all functions of the form

$$\forall x \in \mathcal{X}, \quad K_x : t \mapsto K(x, t) .$$

- ② For every  $x \in \mathcal{X}$  and  $f \in \mathcal{H}$  the **reproducing property** holds:

$$f(x) = \langle f, K_x \rangle_{\mathcal{H}} .$$

If a r.k. exists, then  $\mathcal{H}$  is called a **reproducing kernel Hilbert space (RKHS)**.

# An equivalent definition of RKHS

## Theorem

The Hilbert space  $\mathcal{H} \subset \mathbb{R}^{\mathcal{X}}$  is a RKHS if and only if for any  $x \in \mathcal{X}$ , the mapping:

$$\begin{aligned} F : \mathcal{H} &\rightarrow \mathbb{R} \\ f &\mapsto f(x) \end{aligned}$$

is **continuous**.

# An equivalent definition of RKHS

## Theorem

The Hilbert space  $\mathcal{H} \subset \mathbb{R}^{\mathcal{X}}$  is a RKHS if and only if for any  $x \in \mathcal{X}$ , the mapping:

$$\begin{aligned} F : \mathcal{H} &\rightarrow \mathbb{R} \\ f &\mapsto f(x) \end{aligned}$$

is **continuous**.

## Corollary

**Convergence in a RKHS implies pointwise convergence**, i.e., if  $(f_n)_{n \in \mathbb{N}}$  converges to  $f$  in  $\mathcal{H}$ , then  $(f_n(x))_{n \in \mathbb{N}}$  converges to  $f(x)$  for any  $x \in \mathcal{X}$ .

If  $\mathcal{H}$  is a RKHS then  $f \mapsto f(x)$  is continuous

If a r.k.  $K$  exists, then for any  $(x, f) \in \mathcal{X} \times \mathcal{H}$ :

$$\begin{aligned} |f(x)| &= |\langle f, K_x \rangle_{\mathcal{H}}| \\ &\leq \|f\|_{\mathcal{H}} \cdot \|K_x\|_{\mathcal{H}} \text{ (Cauchy-Schwarz)} \\ &\leq \|f\|_{\mathcal{H}} \cdot K(x, x)^{\frac{1}{2}}, \end{aligned}$$

because  $\|K_x\|_{\mathcal{H}}^2 = \langle K_x, K_x \rangle_{\mathcal{H}} = K(x, x)$ . Therefore  $f \in \mathcal{H} \mapsto f(x) \in \mathbb{R}$  is a continuous linear mapping.  $\square$

## Proof (Converse)

If  $f \mapsto f(x)$  is continuous then  $\mathcal{H}$  is a RKHS

Conversely, let us assume that for any  $x \in \mathcal{X}$  the linear form  $f \in \mathcal{H} \mapsto f(x)$  is continuous.

Then by Riesz representation theorem there (general property of Hilbert spaces) there exists a unique  $g_x \in \mathcal{H}$  such that:

$$f(x) = \langle f, g_x \rangle_{\mathcal{H}}$$

The function  $K(x, y) = g_x(y)$  is then a r.k. for  $\mathcal{H}$ .  $\square$

# Unicity of r.k. and RKHS

## Theorem

- If  $\mathcal{H}$  is a RKHS, then it has a unique r.k.
- Conversely, a function  $K$  can be the r.k. of at most one RKHS.

# Unicity of r.k. and RKHS

## Theorem

- If  $\mathcal{H}$  is a RKHS, then it has a unique r.k.
- Conversely, a function  $K$  can be the r.k. of at most one RKHS.

## Consequence

This shows that we can talk of "the" kernel of a RKHS, or "the" RKHS of a kernel.



# Proof

If a r.k. exists then it is unique

Let  $K$  and  $K'$  be two r.k. of a RKHS  $\mathcal{H}$ . Then for any  $x \in \mathcal{X}$ :

$$\begin{aligned}\|K_x - K'_x\|_{\mathcal{H}}^2 &= \langle K_x - K'_x, K_x - K'_x \rangle_{\mathcal{H}} \\ &= \langle K_x - K'_x, K_x \rangle_{\mathcal{H}} - \langle K_x - K'_x, K'_x \rangle_{\mathcal{H}} \\ &= K_x(x) - K'_x(x) - K_x(x) + K'_x(x) \\ &= 0.\end{aligned}$$

This shows that  $K_x = K'_x$  as functions, i.e.,  $K_x(y) = K'_x(y)$  for any  $y \in \mathcal{X}$ . In other words,  $K=K'$ .  $\square$

# Proof

If a r.k. exists then it is unique

Let  $K$  and  $K'$  be two r.k. of a RKHS  $\mathcal{H}$ . Then for any  $x \in \mathcal{X}$ :

$$\begin{aligned}\|K_x - K'_x\|_{\mathcal{H}}^2 &= \langle K_x - K'_x, K_x - K'_x \rangle_{\mathcal{H}} \\ &= \langle K_x - K'_x, K_x \rangle_{\mathcal{H}} - \langle K_x - K'_x, K'_x \rangle_{\mathcal{H}} \\ &= K_x(x) - K'_x(x) - K_x(x) + K'_x(x) \\ &= 0.\end{aligned}$$

This shows that  $K_x = K'_x$  as functions, i.e.,  $K_x(y) = K'_x(y)$  for any  $y \in \mathcal{X}$ . In other words,  $K=K'$ .  $\square$

The RKHS of a r.k.  $K$  is unique

Left as exercise.

# An important result

## Theorem

A function  $K : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$  is **p.d.** if and only if it is a **r.k.**

# Proof: r.k. $\implies$ p.d.

- ① A r.k. is **symmetric** because, for any  $(x, y) \in \mathcal{X}^2$ :

$$K(x, y) = \langle K_x, K_y \rangle_{\mathcal{H}} = \langle K_y, K_x \rangle_{\mathcal{H}} = K(y, x).$$

- ② It is **p.d.** because for any  $N \in \mathbb{N}, (x_1, x_2, \dots, x_N) \in \mathcal{X}^N$ , and  $(a_1, a_2, \dots, a_N) \in \mathbb{R}^N$ :

$$\begin{aligned} \sum_{i,j=1}^N a_i a_j K(x_i, x_j) &= \sum_{i,j=1}^N a_i a_j \langle K_{x_i}, K_{x_j} \rangle_{\mathcal{H}} \\ &= \left\| \sum_{i=1}^N a_i K_{x_i} \right\|_{\mathcal{H}}^2 \\ &\geq 0. \quad \square \end{aligned}$$

# Proof: p.d. $\implies$ r.k. (1/4)

- Let  $\mathcal{H}_0$  be the vector subspace of  $\mathbb{R}^{\mathcal{X}}$  spanned by the functions  $\{K_x\}_{x \in \mathcal{X}}$ .
- For any  $f, g \in \mathcal{H}_0$ , given by:

$$f = \sum_{i=1}^m a_i K_{x_i}, \quad g = \sum_{j=1}^n b_j K_{y_j},$$

let:

$$\langle f, g \rangle_{\mathcal{H}_0} := \sum_{i,j} a_i b_j K(x_i, y_j).$$

## Proof: p.d. $\implies$ r.k. (2/4)

- $\langle f, g \rangle_{\mathcal{H}_0}$  does not depend on the expansion of  $f$  and  $g$  because:

$$\langle f, g \rangle_{\mathcal{H}_0} = \sum_{i=1}^m a_i g(x_i) = \sum_{j=1}^n b_j f(y_j) .$$

- This also shows that  $\langle \cdot, \cdot \rangle_{\mathcal{H}_0}$  is a symmetric bilinear form.
- This also shows that for any  $x \in \mathcal{X}$  and  $f \in \mathcal{H}_0$ :

$$\langle f, K_x \rangle_{\mathcal{H}_0} = f(x) .$$

## Proof: p.d. $\implies$ r.k. (3/4)

- $K$  is assumed to be p.d., therefore:

$$\|f\|_{\mathcal{H}_0}^2 = \sum_{i,j=1}^m a_i a_j K(x_i, x_j) \geq 0.$$

In particular Cauchy-Schwarz is valid with  $\langle \cdot, \cdot \rangle_{\mathcal{H}_0}$ .

- By Cauchy-Schwarz we deduce that  $\forall x \in \mathcal{X}$ :

$$|f(x)| = |\langle f, K_x \rangle_{\mathcal{H}_0}| \leq \|f\|_{\mathcal{H}_0} K(x, x)^{\frac{1}{2}},$$

therefore  $\|f\|_{\mathcal{H}_0} = 0 \implies f = 0$ .

- $\mathcal{H}_0$  is therefore a **pre-Hilbert space** endowed with the inner product  $\langle \cdot, \cdot \rangle_{\mathcal{H}_0}$ .

## Proof: p.d. $\implies$ r.k. (4/4)

- For any Cauchy sequence  $(f_n)_{n \geq 0}$  in  $(\mathcal{H}_0, \langle \cdot, \cdot \rangle_{\mathcal{H}_0})$ , we note that:

$$\forall (x, m, n) \in \mathcal{X} \times \mathbb{N}^2, \quad |f_m(x) - f_n(x)| \leq \|f_m - f_n\|_{\mathcal{H}_0} \cdot K(x, x)^{\frac{1}{2}}.$$

Therefore for any  $x$  the sequence  $(f_n(x))_{n \geq 0}$  is Cauchy in  $\mathbb{R}$  and has therefore a limit.

- If we add to  $\mathcal{H}_0$  the functions defined as the pointwise limits of Cauchy sequences, then the space becomes complete and is therefore a Hilbert space, with  $K$  as r.k. (up to a few technicalities, left as exercise).  $\square$



# Application: back to Aronszajn's theorem


## Theorem (Aronszajn, 1950)

$K$  is a p.d. kernel on the set  $\mathcal{X}$  *if and only if* there exists a *Hilbert space*  $\mathcal{H}$  and a mapping

$$\Phi : \mathcal{X} \mapsto \mathcal{H} ,$$

such that, for any  $x, x'$  in  $\mathcal{X}$ :

$$K(x, x') = \langle \Phi(x), \Phi(x') \rangle_{\mathcal{H}} .$$




# Proof of Aronzsajn's theorem: p.d. $\implies$ kernel

- If  $K$  is p.d. over a set  $\mathcal{X}$  then it is the r.k. of a Hilbert space  $\mathcal{H} \subset \mathbb{R}^{\mathcal{X}}$ .
- Let the mapping  $\Phi : \mathcal{X} \rightarrow \mathcal{H}$  defined by:

$$\forall x \in \mathcal{X}, \quad \Phi(x) = K_x.$$

- By the reproducing property we have:

$$\forall (x, y) \in \mathcal{X}^2, \quad \langle \Phi(x), \Phi(y) \rangle_{\mathcal{H}} = \langle K_x, K_y \rangle_{\mathcal{H}} = K(x, y). \quad \square$$



# RKHS of the linear kernel

- Let  $\mathcal{X} = \mathbb{R}^d$  and  $K(x, y) = \langle x, y \rangle_{\mathbb{R}^d}$  be the linear kernel
- The corresponding RKHS consists of functions:

$$x \in \mathbb{R}^d \mapsto f(x) = \sum_i a_i \langle x_i, x \rangle_{\mathbb{R}^d} = \langle w, x \rangle_{\mathbb{R}^d} ,$$

with  $w = \sum_i a_i x_i$ .


- The RKHS is therefore the set of **linear forms** endowed with the following inner product:

$$\langle f, g \rangle_{\mathcal{H}_K} = \langle w, v \rangle_{\mathbb{R}^d} ,$$

when  $f(x) = w^\top x$  and  $g(x) = v^\top x$ .

## RKHS of the linear kernel (cont.)

$$\begin{cases} K_{lin}(x, x') &= x^\top x' . \\ f(x) &= w^\top x , \\ \|f\|_{\mathcal{H}} &= \|w\|_2 . \end{cases}$$



## $\ell_2$ -regularized methods in RKHS

$$f_\beta(x) = \beta^\top \Phi(x), \quad \min_{\beta \in \mathbb{R}^p} \left\{ \frac{1}{n} \sum_{i=1}^n \ell(f_\beta(x_i), y_i) + \lambda \|\beta\|_2^2 \right\}$$

is equivalent to

$$\min_{f \in \mathcal{H}} \left\{ \frac{1}{n} \sum_{i=1}^n \ell(f(x_i), y_i) + \lambda \|f\|_{\mathcal{H}}^2 \right\}$$

where  $\mathcal{H}$  is the RKHS of the kernel  $K(x, x') = \Phi(x)^\top \Phi(x')$ .

# Smoothness functional

## A simple inequality

- By Cauchy-Schwarz we have, for any function  $f \in \mathcal{H}$  and any two points  $x, x' \in \mathcal{X}$ :

$$\begin{aligned} |f(x) - f(x')| &= |\langle f, K_x - K_{x'} \rangle_{\mathcal{H}}| \\ &\leq \|f\|_{\mathcal{H}} \times \|K_x - K_{x'}\|_{\mathcal{H}} \\ &= \|f\|_{\mathcal{H}} \times d_K(x, x'). \end{aligned}$$

- The norm of a function in the RKHS controls **how fast** the function varies over  $\mathcal{X}$  with respect to the **geometry defined by the kernel** (Lipschitz with constant  $\|f\|_{\mathcal{H}}$ ).


## Important message

**Small norm  $\implies$  slow variations.**

# Kernels and RKHS : Summary for supervised learning

The goal is to learn a **prediction function**  $f : \mathcal{X} \rightarrow \mathcal{Y}$  given labeled training data  $(x_i, y_i)_{i=1, \dots, n}$  with  $x_i$  in  $\mathcal{X}$ , and  $y_i$  in  $\mathcal{Y}$ :

$$\min_{f \in \mathcal{F}} \underbrace{\frac{1}{n} \sum_{i=1}^n L(y_i, f(x_i))}_{\text{empirical risk, data fit}} + \underbrace{\lambda \Omega(f)}_{\text{regularization}} .$$



# Kernels and RKHS : Summary for supervised learning

The goal is to learn a **prediction function**  $f : \mathcal{X} \rightarrow \mathcal{Y}$  given labeled training data  $(x_i, y_i)_{i=1, \dots, n}$  with  $x_i$  in  $\mathcal{X}$ , and  $y_i$  in  $\mathcal{Y}$ :

$$\min_{f \in \mathcal{F}} \underbrace{\frac{1}{n} \sum_{i=1}^n L(y_i, f(x_i))}_{\text{empirical risk, data fit}} + \underbrace{\lambda \Omega(f)}_{\text{regularization}} .$$

The labels  $y_i$  are in

- $\{-1, +1\}$  for **binary** classification problems.
- $\{1, \dots, K\}$  for **multi-class** classification problems.
- $\mathbb{R}$  for **regression** problems.
- $\mathbb{R}^k$  for **multivariate regression** problems.



# Kernels and RKHS : Summary for supervised learning

The goal is to learn a **prediction function**  $f : \mathcal{X} \rightarrow \mathcal{Y}$  given labeled training data  $(x_i, y_i)_{i=1, \dots, n}$  with  $x_i$  in  $\mathcal{X}$ , and  $y_i$  in  $\mathcal{Y}$ :

$$\min_{f \in \mathcal{F}} \underbrace{\frac{1}{n} \sum_{i=1}^n L(y_i, f(x_i))}_{\text{empirical risk, data fit}} + \underbrace{\lambda \Omega(f)}_{\text{regularization}} .$$

Example with linear models: logistic regression, SVMs, etc.

- assume there exists a linear relation between  $y$  and features  $x$  in  $\mathbb{R}^p$ .
- $f(x) = w^\top x + b$  is parametrized by  $w, b$  in  $\mathbb{R}^{p+1}$ ;
- $L$  is often a **convex** loss function;
- $\Omega(f)$  is often the squared  $\ell_2$ -norm  $\|w\|^2$ .

# Kernels and RKHS : Summary for supervised learning

The goal is to learn a **prediction function**  $f : \mathcal{X} \rightarrow \mathcal{Y}$  given labeled training data  $(x_i, y_i)_{i=1, \dots, n}$  with  $x_i$  in  $\mathcal{X}$ , and  $y_i$  in  $\mathcal{Y}$ :

$$\min_{f \in \mathcal{F}} \underbrace{\frac{1}{n} \sum_{i=1}^n L(y_i, f(x_i))}_{\text{empirical risk, data fit}} + \underbrace{\lambda \Omega(f)}_{\text{regularization}}.$$

## Remark about multilayer neural networks

- The “neural network” space  $\mathcal{F}$  is explicitly parametrized by:

$$f(\mathbf{x}) = \sigma_k(\mathbf{A}_k \sigma_{k-1}(\mathbf{A}_{k-1} \dots \sigma_2(\mathbf{A}_2 \sigma_1(\mathbf{A}_1 \mathbf{x})) \dots)).$$

- Linear operations are either unconstrained (fully connected) or involve parameter sharing (e.g., convolutions).
- Finding the optimal  $\mathbf{A}_1, \mathbf{A}_2, \dots, \mathbf{A}_k$  yields a **non-convex** optimization problem.

# Kernels and RKHS : Summary for supervised learning


A classical kernel formulation for supervised learning

$$\min_{f \in \mathcal{H}} \frac{1}{n} \sum_{i=1}^n L(y_i, f(x_i)) + \lambda \|f\|_{\mathcal{H}}^2.$$

- **map** data  $x$  in  $\mathcal{X}$  to a Hilbert space and work with **linear forms**:

$$\Phi : \mathcal{X} \rightarrow \mathcal{H} \quad \text{and} \quad f(x) = \langle \Phi(x), f \rangle_{\mathcal{H}}.$$

- This is done implicitly with a positive definite kernel  $K(x, x') = \langle \Phi(x), \Phi(x') \rangle$ !




# Kernels and RKHS : Summary for supervised learning

What does it mean to map a data point to a function?

Ex: if  $x, y$  in  $\mathbb{R}$  and  $K(x, y) = e^{-\frac{1}{\sigma^2}(x-y)^2}$  is the Gaussian kernel,

$$\Phi(x) : t \mapsto e^{-\frac{1}{\sigma^2}(x-t)^2}$$

$$\Phi(y) : t \mapsto e^{-\frac{1}{\sigma^2}(y-t)^2}$$



- Data points are mapped to Gaussian functions living in a Hilbert space  $\mathcal{H}$ .
- But  $\mathcal{H}$  is much richer and contains much more than Gaussian functions!
- Prediction functions  $f$  live in  $\mathcal{H}$ :  $f(x) = \langle f, \Phi(x) \rangle$ .

# Kernels and RKHS : Summary

- P.d. kernels can be thought of as **inner product** after embedding the data space  $\mathcal{X}$  in some Hilbert space. As such a p.d. kernel defines a **metric** on  $\mathcal{X}$ .
- A realization of this embedding is the **RKHS**, valid without restriction on the space  $\mathcal{X}$  nor on the kernel.
- The RKHS is a space of functions over  $\mathcal{X}$ . The **norm** of a function in the RKHS is related to its degree of **smoothness** w.r.t. the metric defined by the kernel on  $\mathcal{X}$ .
- $\ell_2$ -regularized learning in the feature space can be formulated in the RKHS

$$\min_{f \in \mathcal{H}} \left\{ \frac{1}{n} \sum_{i=1}^n \ell(f(x_i), y_i) + \lambda \|f\|_{\mathcal{H}}^2 \right\}$$

# Outline

## 1 Learning in high dimension

## 2 Learning with $\ell_2$ regularization

- Ridge regression
- Ridge logistic regression
- Linear hard-margin SVM
- Interlude: quick notes on constrained optimization
- Back to hard-margin SVM
- Soft-margin SVM
- Large-margin classifiers

## 3 Learning with kernels

- Kernel methods
- Positive definite kernels and RKHS
- **Kernel examples**
- Multiple Kernel Learning (MKL)

## 4 Conclusion

# Kernel examples

- **Polynomial** (on  $\mathbb{R}^d$ ):

$$K(x, x') = (x \cdot x' + 1)^d$$

- **Gaussian radial basis function (RBF)** (on  $\mathbb{R}^d$ )

$$K(x, x') = \exp\left(-\frac{\|x - x'\|^2}{2\sigma^2}\right)$$

- **Laplace** kernel (on  $\mathbb{R}$ )

$$K(x, x') = \exp(-\gamma|x - x'|)$$

- **Min** kernel (on  $\mathbb{R}_+$ )

$$K(x, x') = \min(x, x')$$

## Exercise

*Exercise: for each kernel, find a Hilbert space  $\mathcal{H}$  and a mapping  $\Phi : \mathcal{X} \rightarrow \mathcal{H}$  such that  $K(x, x') = \langle \Phi(x), \Phi(x') \rangle$*

# How to choose or make a kernel?

- Design **features**
- Design a **distance** or **similarity** measure
- Design a **regularizer** on  $f$



# Outline

- 1 Learning in high dimension
- 2 Learning with  $\ell_2$  regularization
  - Ridge regression
  - Ridge logistic regression
  - Linear hard-margin SVM
  - Interlude: quick notes on constrained optimization
  - Back to hard-margin SVM
  - Soft-margin SVM
  - Large-margin classifiers
- 3 Learning with kernels
  - Kernel methods
  - Positive definite kernels and RKHS
  - Kernel examples
  - Multiple Kernel Learning (MKL)
- 4 Conclusion

## In one slide...

- Learning in high dimension requires **regularization**, e.g., by  $\ell_2$  penalty for linear methods
- Kernels allow to transform any  $\ell_2$ -regularized linear models into a **nonlinear** model, thanks to the **kernel trick**
- There exists many kernels, which correspond to different **feature spaces** (of finite or infinite dimensions)
- We can **combine** and **learn** kernels, e.g., for integration of heterogeneous data
- Hot research topics
  - **Large-scale** ML with kernels
  - **Deep** kernel methods

## In one slide...

- Learning in high dimension requires **regularization**, e.g., by  $\ell_2$  penalty for linear methods
- Kernels allow to transform any  $\ell_2$ -regularized linear models into a **nonlinear** model, thanks to the **kernel trick**
- There exists many kernels, which correspond to different **feature spaces** (of finite or infinite dimensions)
- We can **combine** and **learn** kernels, e.g., for integration of heterogeneous data
- Hot research topics
  - **Large-scale** ML with kernels
  - **Deep** kernel methods

MURAKOZE